5.3 CoT Prompts

As shown in Table 2, incorporating a guided process improves query quality; however, this comes with trade-offs and requires careful consideration of certain aspects.

While we anticipated that applying CoT would improve query quality and reduce the number of instances required for training, this was not consistently observed, particularly for the TREC-Covid dataset in figure 5. This could be due to several factors:

- The TREC-Covid dataset represents a corpus of 171k documents. Training on a subset of the highest scoring documents seems to decrease the performance of the reranker. Training on higher number of samples, instead of filtering, seems to provide better training signal translated into better downstream performance.
- Although CoT reasoning introduces guided steps for query generation, this can sometimes result in diminished performance. CoT prompts may lead the model to generalize in a way that detracts from the relevance of the generated queries. For instance, in some cases, the reasoning steps cause the model to focus on peripheral concepts rather than extracting keywords or core topics directly relevant to the document.

These results, however, relate to the specific the datasets we test on. Further investigation across a larger pool of corpora is still needed. To answer then the question whether a filtering strategy is needed, the answer seems not evident as of now. Nonetheless, we note that using dynamic prompts instead of a static hardcoded template does improves the downstream performance of the reranker model, as indicated in the top scores for the 4 datasets shown.

Table 4: Sample document from Scifact with its reference query (finding) and queries from the generator models.

	Text	Score
Document	Myelodysplastic syndromes (MDS) are age-dependent stem cell malignancies that...	100
Reference *Finding*	Toll-like receptor (TLR) signaling is involved in the pathogenesis of human MDS.	36.1
Teacher *Finding*	Myeloid-derived suppressor cells (MDSC) are... hematopoiesis. MSDC expansion...	96.3
Student *Finding*	Myeloid-derived suppressor cells (MDSC) are ... of ineffective hematopoiesis.	94.2
CPO *Finding*	Myelodysplastic syndromes (MDS) are age-dependent stem cell malignancies that...	**99.6**

6 Conclusion

In this study, we reproduce and verify the main claims of InPars Toolkit, confirming its claims of providing an end-to-end reproducible pipeline and a clear guide for reproducing InPars-V1, InPars-V2, and partially Promptagator. We also validate the pipeline's plug-and-play capability by integrating LLaMA 3.1 8B as the generator model. Our results closely match the ones reported by the original authors, supporting the toolkit's reliability.

Beyond reproduction, we extend the toolkit in two ways. First, we apply Contrastive Preference Optimization (CPO) to fine-tune the generator LLM to guide the model into generating generally better queries. Second, we apply CoT for dynamic prompt optimization in lieu of fixed templates. This generally improved query quality, though some cases revealed complexities which deserve further investigation.

Overall, our reproductions and extension experiments build upon the versatility of the InPars Toolkit and highlight practical pathways for future research.

References

Hugo Abonizio, Luiz Bonifacio, Vitor Jeronymo, Roberto Lotufo, Jakub Zavrel, and Rodrigo Nogueira. 2023. "Inpars toolkit: A unified and reproducible synthetic data generation pipeline for neural information retrieval." *arXiv preprint arXiv:2307.04601*.

Chris Alberti, Daniel Andor, Emily Pitler, Jacob Devlin, and Michael Collins. 2019. "Synthetic QA corpora generation with roundtrip consistency." *arXiv preprint arXiv:1906.05416*.

Payal Bajaj et al.. 2016. "Ms marco: A human generated machine reading comprehension dataset." *arXiv preprint arXiv:1611.09268*.

Luiz Bonifacio, Hugo Abonizio, Marzieh Fadaee, and Rodrigo Nogueira. 2022. "Inpars: Unsupervised dataset generation for information retrieval." In: *Proceedings of the 45th International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2387–2392.

Leonid Boytsov, Preksha Patel, Vivek Sourabh, Riddhi Nisar, Sayani Kundu, Ramya Ramanathan, and Eric Nyberg. 2023. "Inpars-light: Cost-effective unsupervised training of efficient rankers." *arXiv preprint arXiv:2301.02998*.

Aditi Chaudhary, Karthik Raman, and Michael Bendersky. 2023. *It's All Relative! – A Synthetic Query Generation Approach for Improving Zero-Shot Relevance Prediction*. (2023). https://arxiv.org/abs/2311.07930 arXiv: 2311.07930 [cs.CL].

Zhuyun Dai et al.. 2022. "Promptagator: Few-shot dense retrieval from 8 examples." *arXiv preprint arXiv:2209.11755*.

Aaron Grattafiori et al.. 2024. *The Llama 3 Herd of Models*. (2024). https://arxiv.org/abs/2407.21783 arXiv: 2407.21783 [cs.AI].

Vitor Jeronymo, Luiz Bonifacio, Hugo Abonizio, Marzieh Fadaee, Roberto Lotufo, Jakub Zavrel, and Rodrigo Nogueira. 2023. "Inpars-v2: Large language models as efficient dataset generators for information retrieval." *arXiv preprint arXiv:2301.01820*.

Omar Khattab et al.. 2024. "DSPy: Compiling Declarative Language Model Calls into Self-Improving Pipelines." In.

Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. "Efficient Memory Management for Large Language Model Serving with PagedAttention." In: *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles*.

Patrick Lewis, Yuxiang Wu, Linqing Liu, Pasquale Minervini, Heinrich Küttler, Aleksandra Piktus, Pontus Stenetorp, and Sebastian Riedel. 2021. "PAQ: 65 million probably-asked questions and what you can do with them." *Transactions of the Association for Computational Linguistics*, 9, 1098–1115.

Rodrigo Nogueira, Zhiying Jiang, and Jimmy Lin. 2020. "Document ranking with a pretrained sequence-to-sequence model." *arXiv preprint arXiv:2003.06713*.

R OpenAI. 2023. "Gpt-4 technical report. arxiv 2303.08774." *View in Article*, 2, 5.

Adam Paszke et al.. 2019. *PyTorch: An Imperative Style, High-Performance Deep Learning Library*. (2019). https://arxiv.org/abs/1912.01703 arXiv: 1912.01703 [cs.LG].

Nils Reimers and Iryna Gurevych. Nov. 2019. "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks." In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. Ed. by Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan. Association for Computational Linguistics, Hong Kong, China, (Nov. 2019), 3982–3992. DOI: 10.18653/v1/D19-1410.

David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi. Nov. 2020. "Fact or Fiction: Verifying Scientific Claims." In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Ed. by Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu. Association for Computational Linguistics, Online, (Nov. 2020), 7534–7550. DOI: 10.18653/v1/2020.emnlp-main.609.

Ben Wang and Aran Komatsuzaki. 2021. *GPT-J-6B: A 6 billion parameter autoregressive language model*. (2021).

Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. 2021. "Finetuned language models are zero-shot learners." *arXiv preprint arXiv:2109.01652*.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H Chi, Quoc V Le, and Denny Zhou. 2022. "Chain of Thought Prompting Elicits Reasoning in Large Language Models." *arXiv preprint arXiv:2201.11903*.

Haoran Xu, Amr Sharaf, Yunmo Chen, Weiting Tan, Lingfeng Shen, Benjamin Van Durme, Kenton Murray, and Young Jin Kim. 2024. "Contrastive preference optimization: Pushing the boundaries of llm performance in machine translation." *arXiv preprint arXiv:2401.08417*.